AI for Good and AI Safety in the Development Sector:

A Risk-Governance Framework for NGOs Adopting Generative AI in Grant Writing and Program Management

Author - Nasimul Haque

June 2026

Abstract

Generative artificial intelligence (AI) is increasingly used by non-governmental organizations (NGOs) for grant proposal drafting, donor reporting, monitoring and evaluation (M&E), and day-to-day program management. This adoption promises significant efficiency gains for resource-constrained organizations, but it also introduces risks specific to the development sector: beneficiary data exposure, algorithmic bias in needs assessment, fabricated or inaccurate content in donor-facing documents, erosion of community consent, and over-reliance on automation in decisions that affect vulnerable populations. This paper synthesizes existing AI governance frameworks - including the OECD AI Principles, the NIST AI Risk Management Framework, and humanitarian data protection guidance - and adapts them into a practical, lightweight governance model for small and mid-sized development organizations: the AI Governance and Risk Assessment for Development Organizations (AI-GRADO) framework. The framework comprises five components: purpose screening, data sensitivity tiering, human-in-the-loop checkpoints, bias and equity review, and transparency and disclosure. We illustrate its application through two NGO workflows - grant proposal development and program M&E reporting - and discuss implementation constraints typical of the sector, including limited technical capacity and donor compliance requirements. We conclude that responsible AI adoption in development work is achievable without enterprise-scale resources, provided organizations adopt structured, proportionate safeguards rather than either unmanaged adoption or blanket avoidance.

1. Introduction

The development and humanitarian sector operates under a dual mandate: deliver measurable impact with constrained funding, and uphold a duty of care toward the communities it serves. Generative AI tools - large language models capable of drafting text, summarizing data, and supporting decision-making - have begun to reshape how NGOs operate, particularly in two high-friction areas: grant writing and project management. Grant writing is time-intensive, repetitive across donors, and central to organizational survival; project management involves continuous reporting, indicator tracking, and stakeholder communication, all of which AI can accelerate.

At the same time, the broader AI safety literature has documented failure modes - hallucination, bias amplification, privacy leakage, and automation bias - that carry distinct consequences when applied to populations already in vulnerable circumstances. A fabricated statistic in a donor report can misdirect resources; a biased prioritization algorithm can systematically exclude marginalized groups from aid; and beneficiary data processed through a third-party AI service can violate both organizational policy and the consent under which that data was originally collected.

Despite this, most AI governance frameworks have been written for large enterprises or governments, not for an NGO program director managing a small team with limited technical staff. This paper addresses that gap. It asks: how can a development organization adopt generative AI for grant writing and project management in a way that captures efficiency benefits while remaining accountable to beneficiaries, donors, and regulators? We answer this question by synthesizing relevant literature and proposing a practical governance framework, illustrated with sector-specific use cases.

2. Literature Review

2.1 AI for Good and the Development Sector

The "AI for Good" movement, associated with initiatives such as the International Telecommunication Union's AI for Good platform and the United Nations Sustainable Development Goals (SDGs), frames AI as a potential accelerator for development outcomes in health, education, agriculture, and disaster response. Within this framing, generative AI is increasingly positioned as an operational tool rather than only a frontline service: it supports the administrative backbone of NGOs - proposal writing, translation, data entry, and reporting - that otherwise consumes a disproportionate share of staff time relative to direct programming.

Sector surveys consistently describe two constraints that distinguish NGO AI adoption from corporate adoption: chronic underinvestment in technical infrastructure, and donor compliance regimes that impose strict requirements on data handling and reporting accuracy. Any governance approach for this sector must therefore be lightweight, low-cost, and compatible with existing donor accountability structures rather than introducing a parallel compliance burden.

2.2 AI Safety and Risk Frameworks

General-purpose AI governance has matured around a small number of widely referenced frameworks. The OECD AI Principles establish high-level commitments to transparency, accountability, robustness, and human-centered values. The U.S. National Institute of Standards and Technology (NIST) AI Risk Management Framework offers a process-oriented model organized around four functions - govern, map, measure, and manage - that has been influential across sectors. The EU's risk-tiered approach to AI regulation similarly classifies AI use cases by potential harm, with higher-risk applications subject to stricter obligations.

These frameworks share common safety concerns relevant to NGOs: the risk of automation bias, where users over-trust AI outputs without verification; the risk of data leakage, where sensitive inputs are retained or exposed by third-party AI providers; and the risk of representational harm, where AI systems trained on non-representative data reproduce or amplify existing inequities. None of these frameworks, however, were designed with the specific operating conditions of a small development NGO in mind - conditions that include working with personally identifiable beneficiary data in fragile contexts, multilingual communication, and reporting obligations to multiple donors with differing standards.

2.3 Humanitarian Data Protection and Responsible AI Guidance

Humanitarian-sector-specific guidance, such as data protection handbooks developed by organizations including the International Committee of the Red Cross, and the Sphere humanitarian standards, emphasize the principle of "do no harm" as it applies to data collection and processing. These sources stress informed consent, data minimization, and the heightened sensitivity of data collected from populations affected by conflict, displacement, or disaster. When generative AI tools are introduced into workflows that touch this kind of data, the existing do-no-harm principle extends naturally into AI-specific safeguards: avoiding the input of identifiable beneficiary data into external AI systems without clear data processing agreements, and avoiding AI-generated outputs that could mischaracterize community needs or consent status.

Taken together, the literature suggests three gaps that this paper addresses: (1) a lack of sector-specific, low-resource governance tools; (2) limited guidance distinguishing low-risk AI use (e.g., drafting assistance) from higher-risk use (e.g., automated beneficiary prioritization); and (3) limited attention to donor-facing transparency as a distinct safety dimension, alongside the more commonly discussed dimensions of bias and privacy.

3. Risk Landscape for AI Use in NGO Operations

Based on the literature reviewed and common NGO workflows, five categories of risk are particularly salient for grant writing and project management contexts:

1. Data exposure - beneficiary names, locations, health status, or vulnerability indicators entered into AI tools without verifying data retention and processing terms.
2. Fabrication and inaccuracy - AI-generated statistics, citations, or impact claims in proposals and reports that are plausible but unverified, risking donor trust and compliance violations if discovered.
3. Bias and exclusion - AI-assisted prioritization or needs-ranking tools that reflect biases in training data, systematically underweighting certain regions, languages, or demographic groups.
4. Automation bias and de-skilling - staff over-relying on AI drafts without critical review, gradually eroding institutional expertise in proposal strategy and analysis.
5. Transparency and consent gaps - donors or communities not being informed that AI tools were used in producing reports or communications that represent the organization's voice.

These risks are not uniform in severity. A grant narrative drafted with AI assistance and reviewed by a human carries low residual risk; an AI tool used to rank which communities receive limited aid carries substantially higher risk. This distinction motivates a tiered approach rather than a single blanket policy, which is the basis of the framework proposed in the next section.

4. Proposed Framework: AI-GRADO

We propose the AI Governance and Risk Assessment for Development Organizations (AI-GRADO) framework, designed to be implementable by a program director or small management team without dedicated AI governance staff. The framework consists of five components, summarized in Table 1 and detailed below.

Component	Purpose	Key Question
1. Purpose Screening	Classify the intended AI use case before adoption	Is this task drafting/support, or decision-making about people?
2. Data Sensitivity Tiering	Determine what data may be input into AI tools	Does this input contain identifiable or sensitive beneficiary data?
3. Human-in-the-Loop Checkpoints	Require human verification before external use	Has a qualified staff member reviewed and approved this output?
4. Bias and Equity Review	Check for systematic exclusion in AI-assisted decisions	Could this output disadvantage a specific group or region?
5. Transparency and Disclosure	Maintain donor and community trust	Have we disclosed AI involvement where relevant policy requires it?

Table 1. The five components of the AI-GRADO framework.

4.1 Purpose Screening

Before adopting an AI tool for a given task, the organization classifies the task into one of three tiers: (a) low-risk drafting support (e.g., first-draft proposal narratives, formatting assistance, translation support), (b) moderate-risk analytical support (e.g., summarizing survey data, drafting M&E indicators), or (c) high-risk decision support (e.g., ranking beneficiaries for limited aid, automating eligibility determinations). Tier (c) use cases require senior management sign-off and the full application of components 2 through 5; tier (a) use cases require only basic human review.

4.2 Data Sensitivity Tiering

Inputs to AI tools are classified as public, internal, or sensitive/personally identifiable. Sensitive data - including beneficiary names, exact locations, health information, or protection-related case details - should not be entered into general-purpose AI tools unless the organization has verified the provider's data handling terms and, where required by donor or national law, obtained a data processing agreement. In practice, this often means anonymizing or aggregating data before using it in an AI-assisted summary or report.

4.3 Human-in-the-Loop Checkpoints

No AI-generated content - whether a grant narrative, a donor report, or a program update - should be transmitted externally without review by a staff member with subject-matter knowledge of the program in question. This checkpoint serves two functions: catching factual errors or fabricated details, and preserving institutional voice and accountability, since AI tools cannot be held accountable for representations made to donors or regulators.

4.4 Bias and Equity Review

For any AI-assisted process that influences resource allocation or prioritization - even informally, such as an AI tool suggesting which project sites to highlight in a report - the organization should periodically review outputs for patterns of exclusion (for example, systematic under-representation of a particular language group, geography, or demographic). This need not require technical bias auditing tools; a periodic manual review by program staff comparing AI-suggested outputs against ground-level program data is often sufficient at small-organization scale.

4.5 Transparency and Disclosure

The organization maintains a simple internal record of where AI tools were used in donor-facing materials, and discloses this where donor policy, national regulation, or organizational ethics codes require it. This is increasingly relevant as donors update their own compliance requirements to address AI-assisted reporting; proactive disclosure reduces the risk of trust violations being discovered rather than disclosed.

5. Application: Two NGO Workflows

5.1 Grant Proposal Development

In a typical workflow, a program director uses a generative AI tool to draft sections of a grant proposal based on prior proposals, program data, and donor guidelines. Under AI-GRADO, this is classified as tier (a)/(b): low-to-moderate risk. The data sensitivity check ensures that only aggregated or already-public program statistics are provided as input, rather than raw beneficiary records. The human-in-the-loop checkpoint requires that a staff member with direct program knowledge verify all factual claims, statistics, and impact figures before submission, since AI-generated numbers can be plausible but incorrect. Where the donor's application process requires disclosure of drafting tools used, this is documented per the transparency component.

5.2 Program Monitoring and Evaluation Reporting

AI tools are increasingly used to summarize survey responses, draft narrative sections of M&E reports, or identify patterns across indicator data. This is typically a tier (b) use case, but can shift to tier (c) if AI outputs are used to inform decisions about which communities or individuals receive continued support. In the latter case, AI-GRADO requires senior management sign-off, a documented bias and equity review comparing AI-suggested conclusions against raw field data, and explicit human decision-making authority retained over any resource allocation outcome - the AI tool may inform but not determine such decisions.

6. Discussion

The AI-GRADO framework is intentionally lightweight relative to enterprise AI governance models. This reflects a deliberate trade-off: a framework requiring dedicated compliance staff or specialized technical tooling is unlikely to be adopted by the resource-constrained organizations that make up most of the development sector. Instead, the framework relies on role clarity (who screens, who reviews, who signs off) and proportionality (matching the rigor of safeguards to the actual risk tier of the task), both of which can be implemented through policy documents and staff training rather than new technical infrastructure.

A limitation of this approach is that it depends on consistent staff judgment in classifying tasks correctly into risk tiers, particularly distinguishing moderate-risk analytical support from high-risk decision support. Smaller organizations may also lack the capacity to verify AI providers' data handling terms in detail, which suggests a continued role for donor consortia or sector bodies in negotiating standard data protection terms with AI vendors on behalf of multiple NGOs - reducing the burden on any single organization.

There is also a tension between efficiency gains and the human-in-the-loop requirement: if every AI output requires substantial human review, some of the time savings that motivated adoption in the first place may be eroded. In practice, this tension is best managed by reserving intensive review for higher-stakes outputs (donor reports, factual claims, prioritization decisions) while allowing lighter-touch review for low-stakes drafting (internal notes, first-draft formatting), consistent with the tiered structure proposed above.

7. Limitations of This Paper

This paper presents a conceptual framework synthesized from existing literature and sector practice, rather than results from a controlled empirical study or pilot deployment. The framework has not yet been validated through formal field testing across multiple organizations, and its effectiveness in reducing actual incidents of data exposure, bias, or misinformation remains an empirical question for future research. Future work should pilot AI-GRADO with a small cohort of development organizations, measuring both efficiency outcomes (time saved in proposal and report drafting) and safety outcomes (incidence of factual errors, data handling issues, or bias-related concerns identified through review).

8. Conclusion

Generative AI offers genuine efficiency gains for NGOs facing chronic resource constraints, particularly in grant writing and project management. Realizing these gains safely requires governance that is proportionate to organizational capacity rather than modeled on enterprise compliance regimes. The AI-GRADO framework proposed in this paper - built on purpose screening, data sensitivity tiering, human-in-the-loop checkpoints, bias and equity review, and transparency and disclosure - offers one practical path toward that goal. Responsible AI adoption in the development sector is not a binary choice between unmanaged use and avoidance; it is an exercise in proportionate, role-based safeguards that protect beneficiaries and donor trust while preserving the efficiency benefits that motivated adoption in the first place.

References

International Committee of the Red Cross. Handbook on Data Protection in Humanitarian Action.

National Institute of Standards and Technology. AI Risk Management Framework (AI RMF 1.0).

Organisation for Economic Co-operation and Development. OECD AI Principles.

Sphere Association. The Sphere Handbook: Humanitarian Charter and Minimum Standards in Humanitarian Response.

International Telecommunication Union. AI for Good Global Summit Proceedings.

United Nations. Sustainable Development Goals and the Role of Emerging Technologies.

European Commission. Regulatory Framework Proposal on Artificial Intelligence (EU AI Act).